
INTRODUCTION OF NATURAL LANGUAGE PROCESSING

SHIJU SS



HUMAN LANGUAGE

- The most difficult problem
 - Why ?
- Still not completely solved
- Is data available ?



AMBIGUITY OF LANGUAGE

Classic Example :

Boy saw a man on a hill with a telescope

Look at the dog with one eye

A program

make

HISTORY

- 1950s
 - Early MT: word translation + re-ordering
 - Chomsky's Generative grammar
 - Bar-Hill's argument
- 1960-80s
 - Applications
 - BASEBALL: use NL interface to search in a database on baseball games
 - LUNAR: NL interface to search in Lunar
 - ELIZA: simulation of conversation with a psychoanalyst
 - SHREDLU: use NL to manipulate block world
 - Message understanding: understand a newspaper article on terrorism
 - Machine translation
 - Methods
 - ATN (augmented transition networks): extended context-free grammar
 - Case grammar (agent, object, etc.)
 - DCG – Definite Clause Grammar
 - Dependency grammar: an element depends on another
- 1990s-2010s
 - Statistical methods
 - Speech recognition
 - Question-answering
- 2010s - now
 - Translation
 - Chabot
 - Conversational Bots
 - Question-answering

CLASSICAL STAGES OF NLP

- Morphological Analysis
 - Individual words are analyzed into their components
- Syntactic Analysis
 - Linear sequence of words are transformed into structures that show how the words relate to each other.
- Semantic Analysis
 - Transformation is made from the input text to an internal representation that reflects the meaning – ANALYSE CONTEXT and MEANING
- Pragmatic Analysis
 - To reinterpret what was said to what was actually meant
- Discourse Analysis
 - Resolving references between sentences. (It. He, they etc.)

TOKENIZATION AND SENTENCE SEGMENTATION

- Cells of a language.
- Same words occur many times in a sentence
- Is word count a relevant feature ?
- How many tokens in “India is my country”
- Words may represent the object which we are looking for.
 - “Total budget for this year is approximately 25000 USD.
- After tokenization need stemming in some analysis.
 - What is stemming?
 - Life -> life
 - life -> life
 - Lives -> life

PART OF SPEECH

- Grammatical tagging
- Word category disambiguation.
- Based on definition and context
 - Make cake
 - Make of a car
- Approach :-
 - mid:1980s – HMM
 - Dynamic Programming – similar to Viterbi
 - Mostly variations of HMM.

CHUNKING

- Identifies phrase level constituents in sentences.
 - Noun phrases – “Prime minister Modi”
 - Verb Phrases – “playing cricket”
- Phrase extraction
- Detecting the mentions of key entities or other relevant information.
- Uses POS tag information along with words
- Most relevant in Named Entity Recognition

RULES VS. STATISTICS

- Rules and categories do not fit a sentence equally
 - Some are more likely in a language than others
 - E.g.
 - hardcopy: noun or verb?
 - $P(N \mid \text{hardcopy}) \gg P(V \mid \text{hardcopy})$
 - the training ...
 - $P(N \mid \text{training, Det}) > P(V \mid \text{training, Det})$
- Idea: use statistics to help

STATISTICAL ANALYSIS TO HELP SOLVE AMBIGUITY

- Choose the most likely solution

$$\text{solution}^* = \operatorname{argmax}_{\text{solution}} P(\text{solution} \mid \text{word}, \text{context})$$

$$\text{e.g. } \operatorname{argmax}_{\text{cat}} P(\text{cat} \mid \text{word}, \text{context})$$

$$\operatorname{argmax}_{\text{sem}} P(\text{sem} \mid \text{word}, \text{context})$$

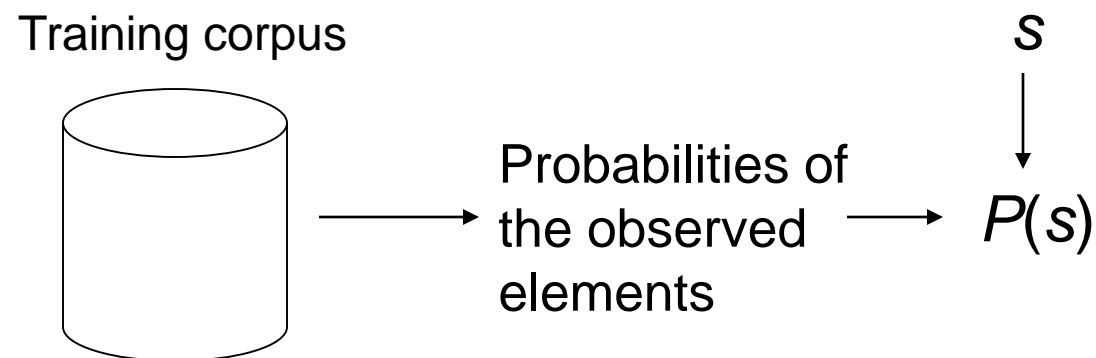
Context varies largely (precedent work, following word, category of the precedent word, ...)

- How to obtain $P(\text{solution} \mid \text{word}, \text{context})$?
 - Training corpus

STATISTICAL LANGUAGE MODELING

Goal: create a statistical model so that one can calculate the probability of a sequence of tokens $s = w_1, w_2, \dots, w_n$ in a language.

- General approach:



PROB. OF A SEQUENCE OF WORDS

$$\begin{aligned} P(s) &= P(w_1, w_2, \dots, w_n) \\ &= P(w_1)P(w_2 | w_1) \dots P(w_n | w_{1,n-1}) \\ &= \prod_{i=1}^n P(w_i | h_i) \end{aligned}$$

Elements to be estimated:

$$P(w_i | h_i) = \frac{P(h_i w_i)}{P(h_i)}$$

- If h_i is too long, one cannot observe (h_i, w_i) in the training corpus, and (h_i, w_i) is hard to generalize
- Solution: limit the length of h_i

N-GRAMS

$$P(s) = \prod_{i=1}^n P(w_i)$$

$$P(s) = \prod_{i=1}^n P(w_i \mid w_{i-1})$$

$$P(s) = \prod_{i=1}^n P(w_i \mid w_{i-2}w_{i-1})$$

JACCARD COEFFICIENT

- A commonly used measure of overlap of two sets A and B
- $\text{jaccard}(A,B) = |A \cap B| / |A \cup B|$
- $\text{jaccard}(A,A) = 1$
- $\text{jaccard}(A,B) = 0$ if $A \cap B = 0$
- A and B don't have to be the same size.
- Always assigns a number between 0 and 1.

JACCARD COEFFICIENT: SCORING EXAMPLE

- What is the query-document match score that the Jaccard coefficient computes for each of the two documents below?
- Query: *ides of march*
- Document 1: *caesar died in march*
- Document 2: *the long march*

ISSUES WITH JACCARD FOR SCORING

- It doesn't consider *term frequency* (how many times a term occurs in a document)
- Rare terms in a collection are more informative than frequent terms. Jaccard doesn't consider this information
- We need a more sophisticated way of normalizing for length

$$|A \cap B| / \sqrt{|A \cup B|}$$

BINARY TERM-DOCUMENT INCIDENCE MATRIX

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

Each document is represented by a binary vector $\in \{0,1\}^{|V|}$

TERM-DOCUMENT COUNT MATRICES

- Consider the number of occurrences of a term in a document:
 - Each document is a count vector in \mathbb{N}^v : a column below

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	157	73	0	0	0	0
Brutus	4	157	0	1	0	0
Caesar	232	227	0	2	1	1
Calpurnia	0	10	0	0	0	0
Cleopatra	57	0	0	0	0	0
mercy	2	0	3	5	5	1
worser	2	0	1	1	1	0

BAG OF WORDS MODEL

- Vector representation doesn't consider the ordering of words in a document
- *John is quicker than Mary* and *Mary is quicker than John* have the same vectors
- This is called the bag of words model.
- In a sense, this is a step back: The positional index was able to distinguish these two documents.
- For now: bag of words model

TERM FREQUENCY TF

- The term frequency $tf_{t,d}$ of term t in document d is defined as the number of times that t occurs in d .
- We want to use tf when computing query-document match scores. But how?
- Raw term frequency is not what we want:
 - A document with 10 occurrences of the term is more relevant than a document with 1 occurrence of the term.
 - But not 10 times more relevant.
- Relevance does not increase proportionally with term frequency.

LOG-FREQUENCY WEIGHTING

- The log frequency weight of term t in d is

$$w_{t,d} = \begin{cases} 1 + \log_{10} \text{tf}_{t,d}, & \text{if } \text{tf}_{t,d} > 0 \\ 0, & \text{otherwise} \end{cases}$$

- $0 \rightarrow 0, 1 \rightarrow 1, 2 \rightarrow 1.3, 10 \rightarrow 2, 1000 \rightarrow 4$, etc.
- Score for a document-query pair: sum over terms t in both q and d :

- score
$$= \sum_{t \in q \cap d} (1 + \log \text{tf}_{t,d})$$

- The score is 0 if none of the query terms is present in the document.

DOCUMENT FREQUENCY

- Rare terms are more informative than frequent terms
 - Recall stop words
- Consider a term in the query that is rare in the collection (e.g., *arachnocentric*)
- A document containing this term is very likely to be relevant to the query *arachnocentric*
- → We want a high weight for rare terms like *arachnocentric*.

DOCUMENT FREQUENCY, CONTINUED

- Frequent terms are less informative than rare terms
- Consider a query term that is frequent in the collection (e.g., *high*, *increase*, *line*)
- A document containing such a term is more likely to be relevant than a document that doesn't
- But it's not a sure indicator of relevance.
- → For frequent terms, we want high positive weights for words like *high*, *increase*, and *line*
- But lower weights than for rare terms.
- We will use document frequency (df) to capture this.

IDF WEIGHT

- df_t is the document frequency of t : the number of documents that contain t
 - df_t is an inverse measure of the informativeness of t
 - $df_t \leq N$
- We define the idf (inverse document frequency) of t by

$$idf_t = \log_{10} (N/df_t)$$

- We use $\log (N/df_t)$ instead of N/df_t to “dampen” the effect of idf.

TF-IDF WEIGHTING

- The tf-idf weight of a term is the product of its tf weight and its idf weight.

$$w_{t,d} = \log(1 + \text{tf}_{t,d}) \times \log_{10}(N / \text{df}_t)$$

- Best known weighting scheme in information retrieval
 - Note: the “-” in tf-idf is a hyphen, not a minus sign!
 - Alternative names: tf.idf, tf x idf
- Increases with the number of occurrences within a document
- Increases with the rarity of the term in the collection

SCORE FOR A DOCUMENT GIVEN A QUERY

$$\text{Score}(q, d) = \sum_{t \in q \cap d} \text{tf.idf}_{t,d}$$

- There are many variants
 - How “tf” is computed (with/without logs)
 - Whether the terms in the query are also weighted
 - ...

BINARY → COUNT → WEIGHT MATRIX

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	5.25	3.18	0	0	0	0.35
Brutus	1.21	6.1	0	1	0	0
Caesar	8.59	2.54	0	1.51	0.25	0
Calpurnia	0	1.54	0	0	0	0
Cleopatra	2.85	0	0	0	0	0
mercy	1.51	0	1.9	0.12	5.25	0.88
worser	1.37	0	0.11	4.15	0.25	1.95

Each document is now represented by a real-valued vector of tf-idf weights $\in \mathbb{R}^{|V|}$

FEATURE ENGINEERING

- Conventional hand crafted features → Representation Learning
- Representation Learning
 - We deal with wide variety of data
 - Signals, Images, Text etc.
 - Which is the best way of representing the data
 - How to represent the data, so that machine can learn from it
 - Def : A bunch of **techniques** which **automatically** detect / **learn** the **representation** for feature selection from natural / raw data.

REPRESENTATIONS IN NLP

- Word Representation
- Sentence Representation
- Document Representation

WORD REPRESENTATIONS

Traditional Method - Bag of Words Model

- Uses one hot encoding
- Each word in the vocabulary is represented by one bit position in a HUGE vector.
- For example, if we have a vocabulary of 10000 words, and “Hello” is the 4th word in the dictionary, it would be represented by: 0 0 0 1 0 0 0 0 0 0
- Context information is not utilized

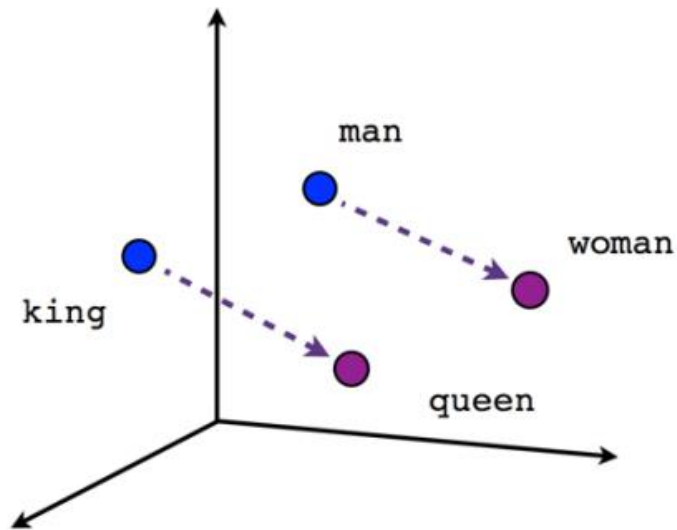
Word Embeddings

- Stores each word in as a point in space, where it is represented by a vector of fixed number of dimensions (generally 300)
- Unsupervised, built just by reading huge corpus
- For example, “Hello” might be represented as :
[0.4, -0.11, 0.55, 0.3 ... 0.1, 0.02]
- Dimensions are basically projections along different axes, more of a mathematical concept.

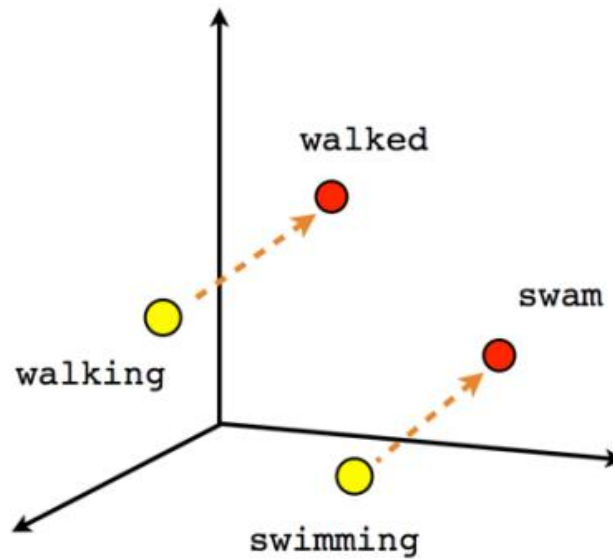
THE POWER OF WORD VECTORS

- They provide a fresh perspective to ***almost ALL*** problems in NLP.
- Technological Improvement
 - Rise of deep learning since 2006 (Big Data + GPUs + Work done by Andrew Ng, Yoshua Bengio, Yann Lecun and Geoff Hinton)
 - Application of Deep Learning to NLP – led by Yoshua Bengio, Christopher Manning, Richard Socher, Tomas Mikalov
- The need for unsupervised learning . (Supervised learning tends to be excessively dependant on hand-labelled data and often does not scale)

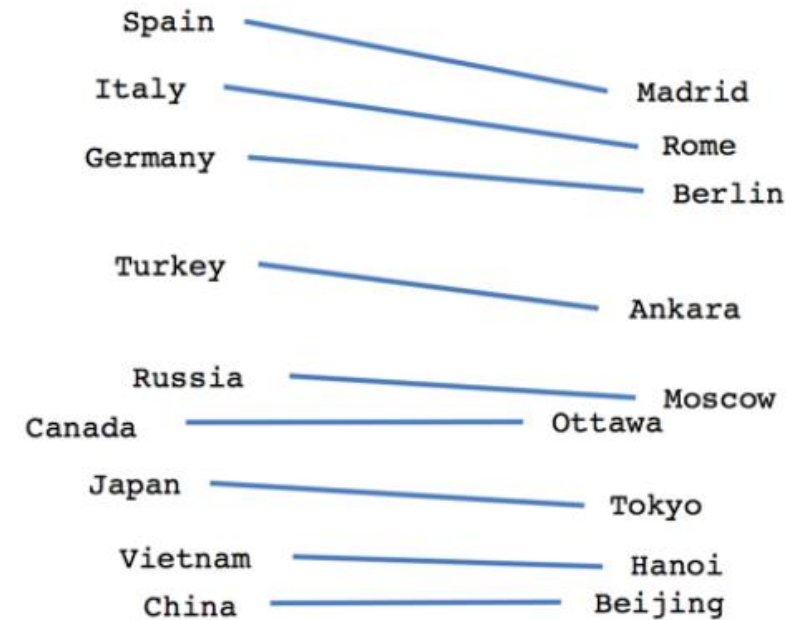
EXAMPLES



Male-Female



Verb tense



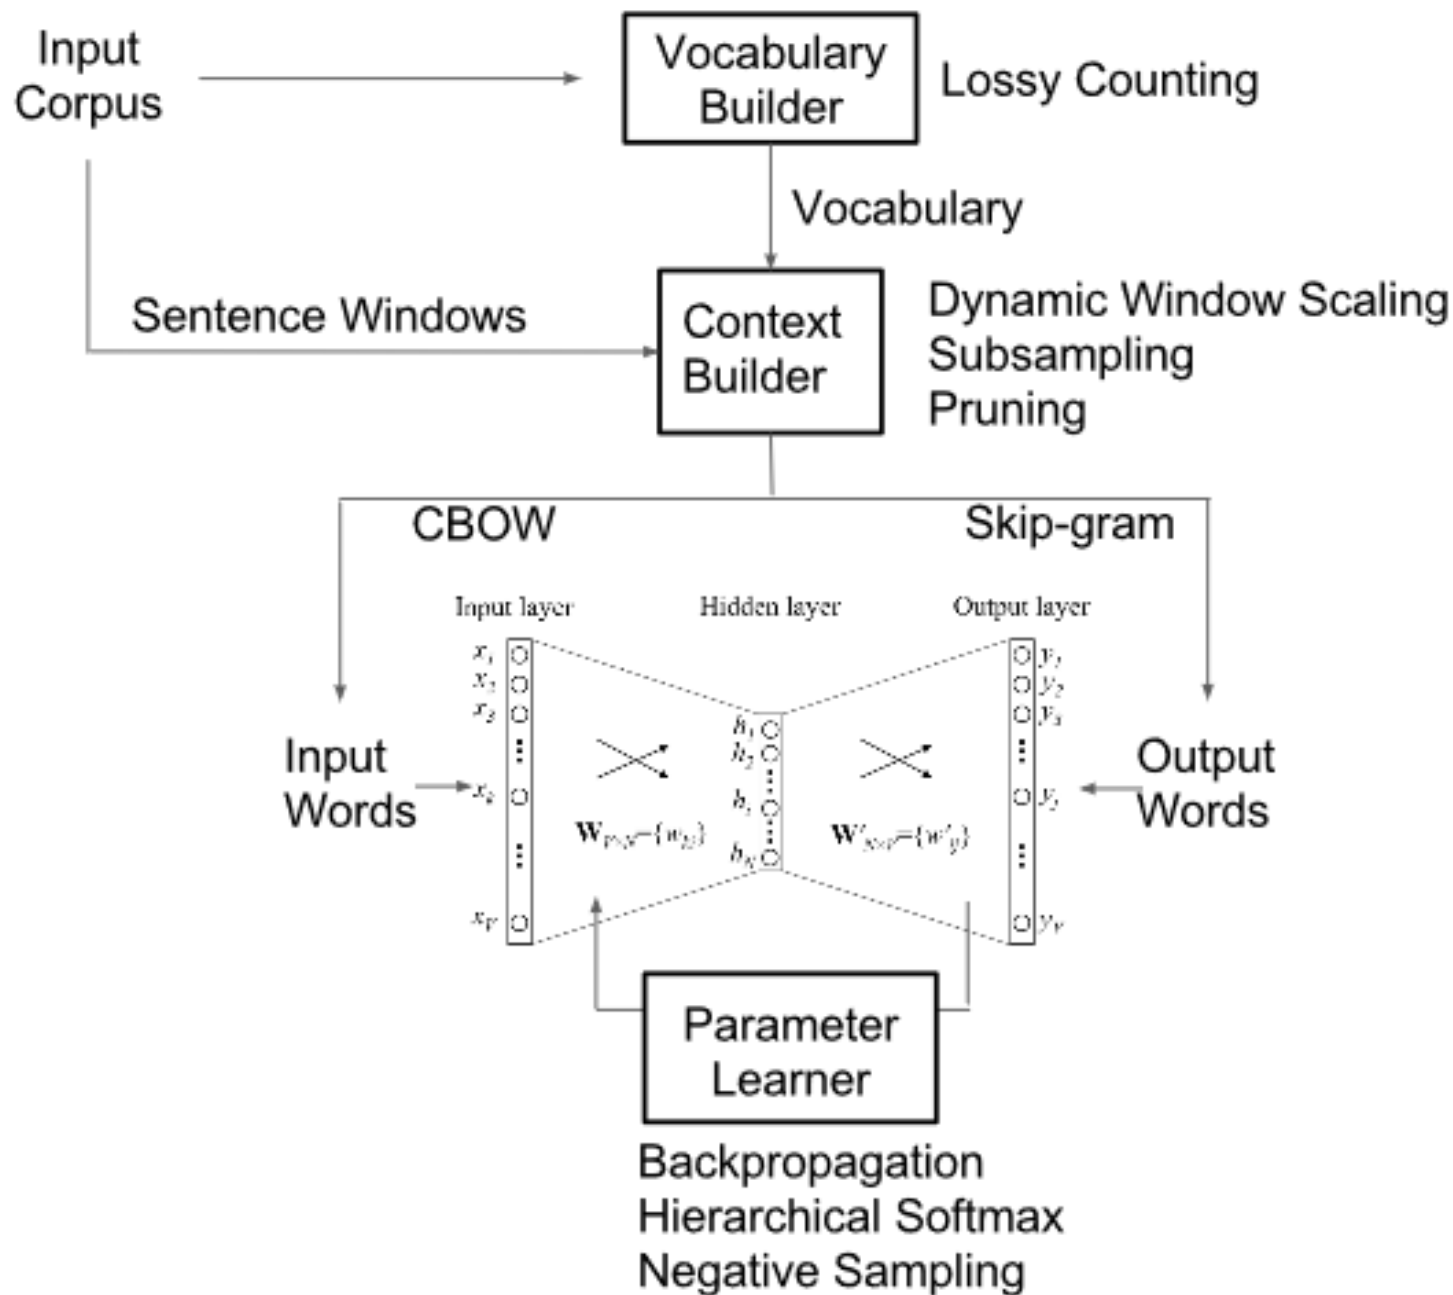
Country-Capital

$$\text{vector[Queen]} = \text{vector[King]} - \text{vector[Man]} + \text{vector[Woman]}$$

ARCHITECTURE

$$\frac{1}{T} \sum_{t=1}^T \sum_{-c \leq j \leq c, j \neq 0} \log p(w_{t+j} | w_t)$$

$$p(w_O | w_I) = \frac{\exp(v'_{w_O} \top v_{w_I})}{\sum_{w=1}^W \exp(v'_w \top v_{w_I})}$$



TRAINING – SKIP GRAM

Sentences

The big brown bear is sitting in a chair.



(the, big)
(the, brown)

The big brown bear is sitting in a chair.



(big, the)
(big, brown)
(big, bear)

The big brown bear is sitting in a chair.



(brown, the)
(brown, big)
(brown, bear)
(brown, is)

The big brown bear is sitting in a chair.



(bear, big)
(bear, brown)
(bear, is)
(bear, sitting)

APPLICATIONS

Word Similarity

Classic Methods : Edit Distance, WordNet, Porter's Stemmer, Lemmatization using dictionaries

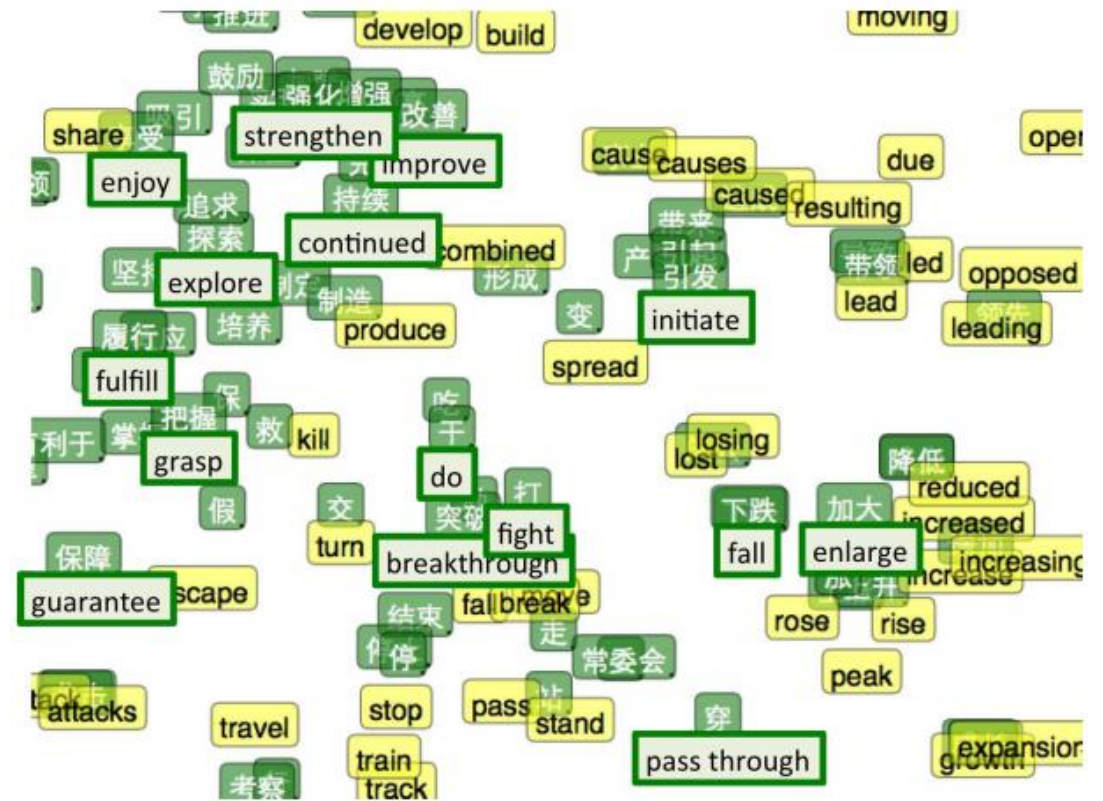
- Easily identifies similar words and synonyms since they occur in similar contexts
- Stemming (thought -> think)
- Inflections, Tense forms
- eg. *Think, thought, ponder, pondering,*
- eg. *Plane, Aircraft, Flight*

NAMED ENTITY RECOGNITION

- Identifies and classifies strings of characters representing proper nouns.
- Useful for filtering documents
- Linguistic grammar based models
 - Months of work by linguistic expert
 - High precision
- Machine learning
 - Supervised as well semi-supervised
 - Require large amount of annotated data
 - Conditional Random Fields

MACHINE TRANSLATION

Classic Methods : Rule-based machine translation, morphological transformation



LANGUAGE TRANSLATION

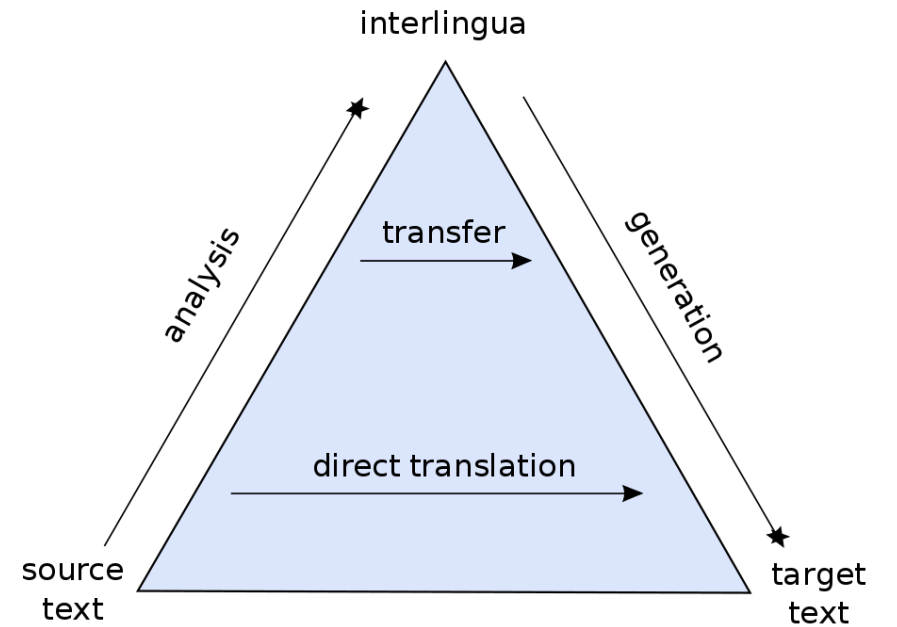
- Languages differs in
 - Script
 - Vocabulary
 - Grammar

These difference can be considered as a measure of language distance

Machine Translation Problem

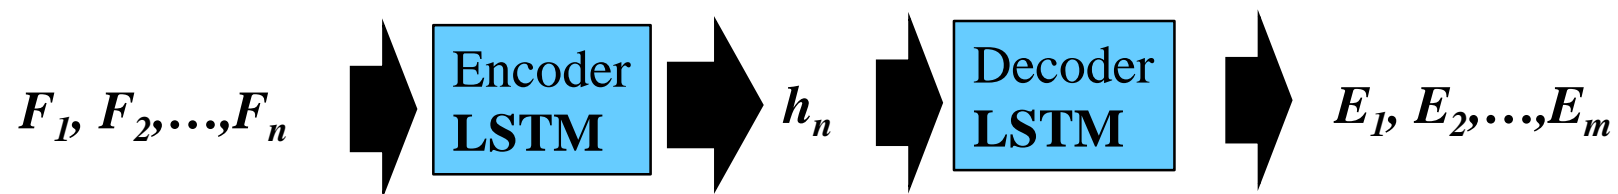
Language Encoding -> representation and analysis

Language Decoding -> synthesis



NEURAL MACHINE TRANSLATION (NMT)

- Encoder/Decoder framework maps sentence in source language to a "deep vector" then another LSTM maps this vector to a sentence in the target language

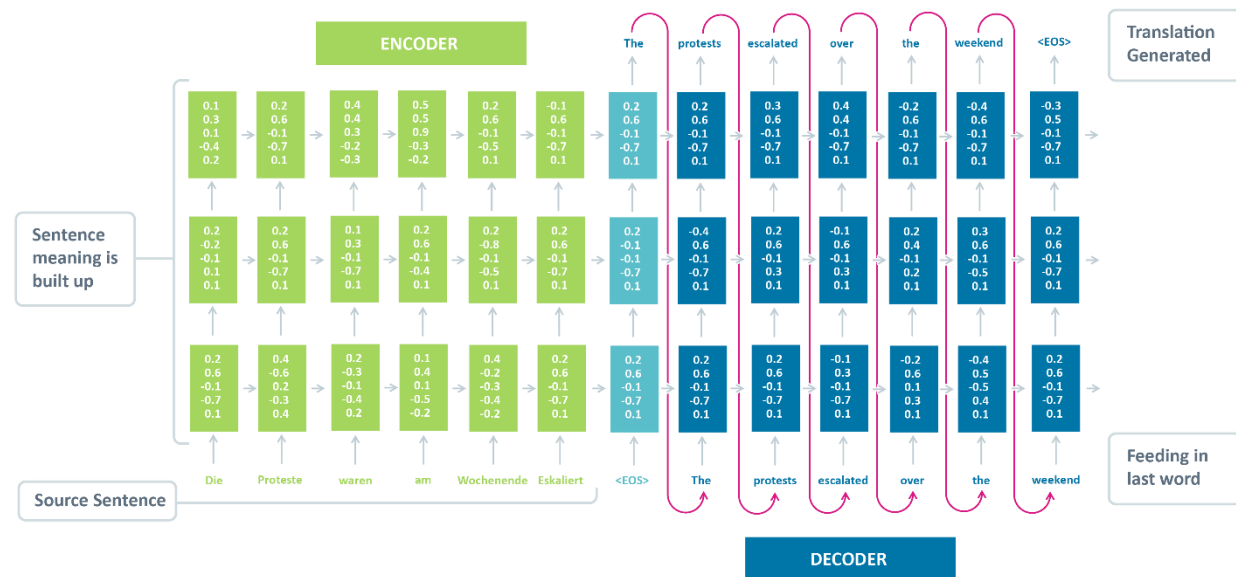


Train model "end to end" on sentence-aligned parallel corpus.

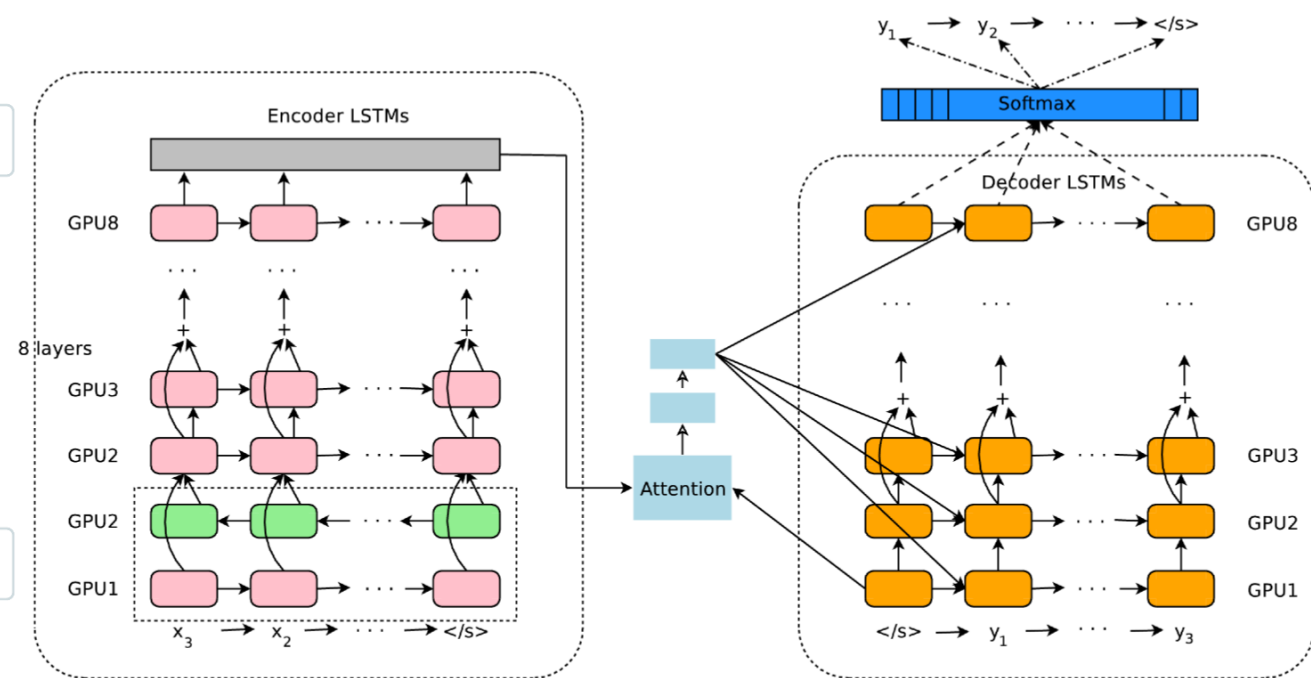
NMT

■ Without Attention

A Recurrent Neural Network for Machine Translation



■ With Attention



CHATBOT

- Chat Bot: A computer program that can ***talk*** to humans in natural language!
- Uses Artificial Intelligence Markup Language (AIML) to represent knowledge.
- Can replace a human for monotonous jobs of answering queries, e.g. E-help desk.

ELIZA

- Eliza – the first chat bot made by *Joseph Weizenbaum*.
- Eliza Effect
 - tendency of humans to attach associations to terms from prior experience.
- Working of Eliza is based on
 - Knowledge Representation
 - Pattern Recognition
 - Substitution of key words into known phrases.

ELIZA

- Looks for certain patterns of words in the user's input.
- Replies with pre-determined output, if the pattern is matched.
- Needs to have an idea of what the user will chat
- Has suitable responses defined in the AIML file

AIML WORKING

- Contains patterns that does not have wildcards “*” or “_”.

- Example:

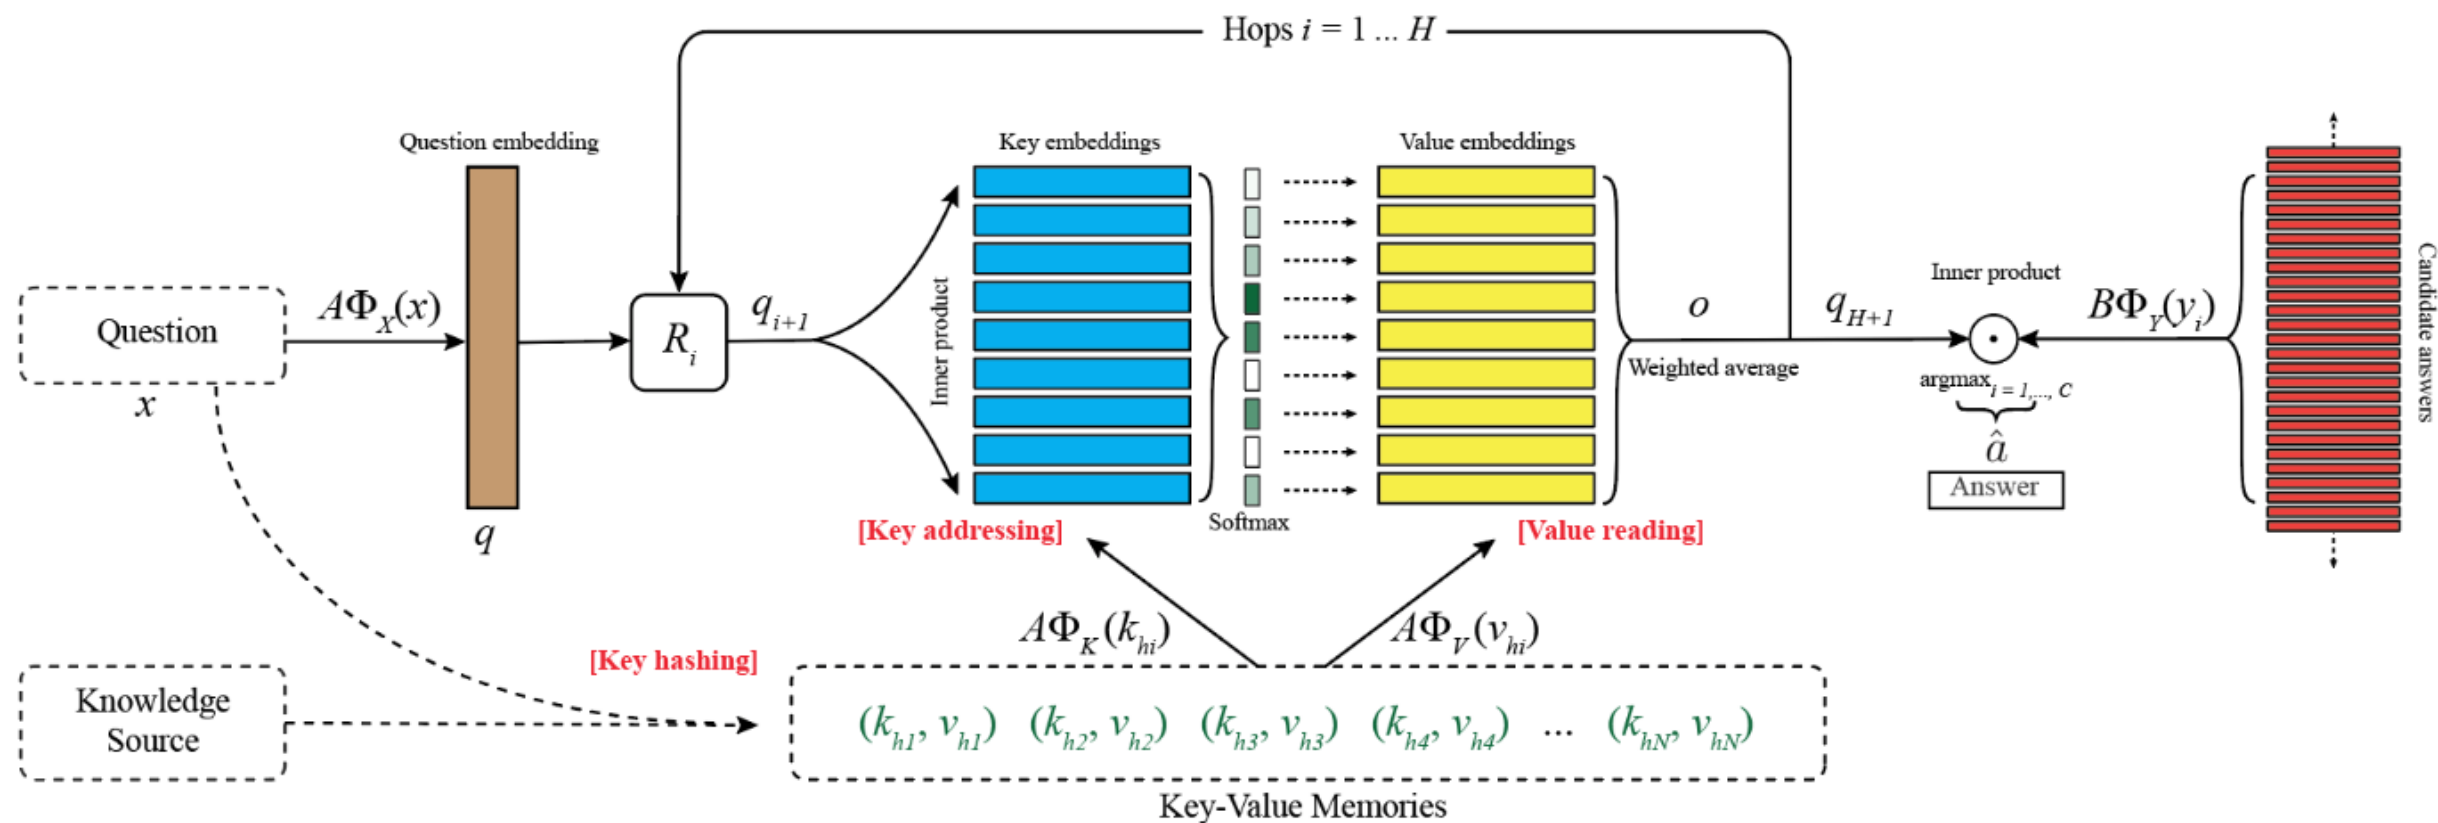
```
<category>  
  <pattern>10 DOLLARS</pattern>  
  <template> wow, what a cheap </template>  
</category>
```

- Conversation:

User: This watch is for 10 dollars

Chat Bot: Wow, what a cheap watch!

MEMORY NETWORKS FOR QA



CONCLUSION

- Started with one hot encoding,
- RF-IDF
- Word2Vec
- Encoder-Decoder
- Memory Networks & Neural Turing Machine
- And goes on.



Questions ?